

Query Language and Groupings

Aggregate Operators and SQL Clauses

Computer Science - IIS Fermi Sacconi Ceci

Ascoli Piceno

All rights reserved

Table of Contents



- Aggregate Operators
- GROUP BY Clause
- HAVING Clause
- Result Limitation

SQL Extension

Compared to relational algebra, the most important extension introduced by SQL is that of **aggregate operators**.

Main characteristics:


- Aggregate multiple table rows
- Perform operations at *relation* level, not tuple level
- Return a **single aggregate value**
- Do not select a subset of rows, but calculate values

Aggregate Operators

Functionality

Aggregation operations allow you to:

- **Group** tuples
- Perform **specific calculations** (sums, counts, statistics)
- **Divide** the table into subsets
- Group tuples with **same values** for specific attributes




Aggregate Operators

Fundamental difference

Aggregation operators, unlike mathematical operators that act only on the tuple being processed:

Can be applied to multiple tuples of the table and are characterized by:

- Returning a **unique value**
- Operating on a **group of values**
- Working on values that form a **column**



Aggregate Operators

Main aggregation operators

`AVG(field)` Calculates the **arithmetic mean**

`COUNT(expr|*)` **Counts** rows

`MAX(expr)` Calculates the **maximum** value

`MIN(expr)` Calculates the **minimum** value

`SUM(field)` Calculates the **total sum**

`STDDEV(field)` Calculates the **standard deviation**

`SELECT AVG(price)` → 234.56

Aggregate Operators

GROUP BY Clause

Specifies which fields to perform groupings on.

Important

Aggregate operators can only appear after these clauses:

- SELECT
- HAVING

```
SELECT category, COUNT(*) AS number  
FROM products  
GROUP BY category;
```

The HAVING Clause

How it works

The HAVING clause allows you to insert a constraint on the data resulting from the GROUP BY grouping operation.

Difference with WHERE:

- WHERE operates on **database fields** (before grouping)
- HAVING operates on **fields resulting from groupings** (after)

```
SELECT department, AVG(salary) AS average
FROM employees
GROUP BY department
HAVING AVG(salary) > 30000;
```


COUNT Operator

Syntax

```
COUNT(expression | *)
```

```
SELECT COUNT(*) AS total_cars  
FROM cars;
```

| total_cars |
|------------|
| 12 |

Note: COUNT(*) counts all rows, including duplicates.

COUNT Operator

Result

The result is **12**: all records that make up the table and are not NULL are counted, regardless of their content.

Important

COUNT is the only aggregation operator that considers null values in the calculation.

To count distinct values:

```
SELECT COUNT(DISTINCT brand) AS  
        different_brands  
FROM cars;
```


| different_brands |
|------------------|
| 8 |

AVG Operator

Arithmetic Mean

Calculates the average of values in a numeric column.

```
SELECT AVG(price) AS average_price  
FROM cars;
```




MIN and MAX Operators

How they work

MAX and MIN return the highest and lowest values contained within a column.

Behavior by data type:

- **Numeric fields:** returns the numeric value
- **Text fields:** identifies the field that, in alphabetical order, is last (MAX) or first (MIN)
- Ordering according to the **ASCII** code table
- **Lexicographic** ordering



MIN and MAX Operators

Mathematical expressions

It is possible to use mathematical expressions as arguments for the MIN and MAX operators.

```
SELECT MAX(price * 1.22) AS max_price_vat  
FROM cars;
```

```
SELECT MIN(price - discount) AS  
       min_price_discounted  
FROM cars;
```

Price \longrightarrow * 1.22 \longrightarrow MAX()

MIN and MAX Operators

```
SELECT MIN(salary) AS minimum_salary  
FROM employees;
```

| Name | Salary |
|----------------|--------------|
| Mario Rossi | 28000 |
| Laura Bianchi | 32000 |
| Giuseppe Verdi | 25000 |
| Anna Neri | 35000 |

| minimum_salary |
|----------------|
| 25000 |

MIN and MAX Operators

```
SELECT MAX(salary) AS maximum_salary  
FROM employees;
```

| Name | Salary |
|----------------|--------------|
| Mario Rossi | 28000 |
| Laura Bianchi | 32000 |
| Giuseppe Verdi | 25000 |
| Anna Neri | 35000 |

| maximum_salary |
|----------------|
| 35000 |

MIN and MAX Operators

```
SELECT MAX(surname) AS max_surname  
FROM employees;
```

| Surname |
|--------------|
| Bianchi |
| Neri |
| Rossi |
| Verdi |

| max_surname |
|--------------|
| Verdi |

Alphabetical order: B | N | R | V

MIN and MAX Operators

```
SELECT MIN(surname) AS min_surname  
FROM employees;
```

| Surname |
|---------|
| Bianchi |
| Neri |
| Rossi |
| Verdi |

| min_surname |
|-------------|
| Bianchi |

Alphabetical order: B | N | R | V

MIN and MAX Operators

Complex expressions

Let's see how to apply arithmetic operators in the query. To know the difference between minimum and maximum employee salaries, expressed as a percentage:

```
SELECT  
  (MAX(salary) - MIN(salary)) /  
  MAX(salary) * 100 AS percentage_diff  
FROM employees;
```

| percentage_diff |
|-----------------|
| 28.57 |


The difference between maximum and minimum salary is 28.57%

SUM Operator

Total sum

Calculates the sum of all values in a numeric column.

```
SELECT SUM(salary) AS total_salaries  
FROM employees;
```



SUM Operator

```
SELECT  
SUM(amount) AS total_sales,  
SUM(quantity) AS total_pieces  
FROM sales  
WHERE year = 2024;
```

| total_sales | total_pieces |
|-------------|--------------|
| 158750.50 | 1243 |


Note

SUM ignores NULL values in the calculation.

AVG and STDDEV Operators

```
SELECT  
  AVG(salary) AS average_salary,  
  STDDEV(salary) AS standard_deviation  
FROM employees;
```

| average_salary | standard_deviation |
|----------------|--------------------|
| 30000.00 | 4082.48 |




The GROUP BY Clause

Function

The GROUP BY clause is used to group and process uniformly different rows that have equal values in a specific column in the source table.

Syntax:

```
SELECT column1, aggregate_function(column2)
FROM table
WHERE condition
GROUP BY column1;
```




The GROUP BY Clause

Execution order

Taking into account the following considerations, let's clarify the priorities during the execution of a selection with grouping.

Processing phases:


- 1 Execution of the JOIN, if present
- 2 Selection on tuples with the condition after WHERE



The GROUP BY Clause

Processing phases (continued):

- ③ Grouping based on fields specified after GROUP BY
- ④ List of values specified after SELECT (target list)
- ⑤ Selection of tuples that satisfy the condition after HAVING



The GROUP BY Clause

```
SELECT department, COUNT(*) AS  
    number_employees  
FROM employees  
GROUP BY department;
```

Original table:


| Name | Department |
|----------|------------|
| Mario | Sales |
| Laura | IT |
| Giuseppe | Sales |
| Anna | IT |
| Carlo | Marketing |

Result:

| department | number_employees |
|------------|------------------|
| Sales | 2 |
| IT | 2 |
| Marketing | 1 |

The GROUP BY Clause

```
SELECT category, AVG(price) AS average_price  
FROM products  
GROUP BY category;
```



The GROUP BY Clause

```
SELECT
year,
month,
SUM(sales) AS total_sales
FROM monthly_sales
GROUP BY year, month
ORDER BY year, month;
```

| year | month | total_sales |
|------|-------|-------------|
| 2023 | 1 | 45000 |
| 2023 | 2 | 52000 |
| 2024 | 1 | 48000 |
| 2024 | 2 | 55000 |

Grouping on **multiple columns**

The GROUP BY Clause

```
SELECT
city,
COUNT(*) AS number_customers,
SUM(orders) AS total_orders
FROM customers
GROUP BY city;
```

| city | number_customers | total_orders |
|--------|------------------|--------------|
| Rome | 15 | 234 |
| Milan | 22 | 387 |
| Naples | 8 | 145 |
| Turin | 12 | 198 |

The GROUP BY Clause

```
SELECT  
brand,  
model,  
COUNT(*) AS units_sold,  
AVG(price) AS average_price  
FROM cars_sold  
GROUP BY brand, model  
ORDER BY units_sold DESC;
```

| brand | model | units | price |
|------------|--------|-------|-------|
| Fiat | Panda | 45 | 12500 |
| Volkswagen | Golf | 38 | 22000 |
| Ford | Fiesta | 32 | 15000 |

The GROUP BY Clause

```
SELECT
YEAR(sale_date) AS year,
MONTH(sale_date) AS month,
COUNT(*) AS number_transactions,
SUM(amount) AS revenue
FROM transactions
GROUP BY YEAR(sale_date), MONTH(sale_date);
```

| year | month | transactions | revenue |
|------|-------|--------------|----------|
| 2024 | 10 | 152 | 45200.00 |
| 2024 | 11 | 178 | 52800.00 |
| 2024 | 12 | 201 | 68500.00 |

Grouping with **functions** in columns

The HAVING Clause

```
SELECT department, AVG(salary) AS  
       average_salary  
FROM employees  
GROUP BY department  
HAVING AVG(salary) > 30000;
```

All groups:

| department | average_salary |
|------------|----------------|
| Sales | 28000 |
| IT | 35000 |
| Marketing | 27000 |
| Production | 32000 |

Result with HAVING:

| department | average_salary |
|------------|----------------|
| IT | 35000 |
| Production | 32000 |

The HAVING Clause

```
SELECT category, COUNT(*) AS number_products
FROM products
GROUP BY category
HAVING COUNT(*) >= 5;
```


| category | number_products |
|-------------|-----------------|
| Electronics | 12 |
| Clothing | 8 |
| Home | 6 |

Note

HAVING filters groups **after** aggregation, while WHERE filters rows **before**.

The HAVING Clause

```
SELECT
customer_id,
SUM(amount) AS total_purchases
FROM orders
WHERE year = 2024
GROUP BY customer_id
HAVING SUM(amount) > 10000;
```



The HAVING Clause

```
SELECT
city,
COUNT(*) AS number_orders,
AVG(amount) AS average_amount
FROM orders
GROUP BY city
HAVING COUNT(*) > 10 AND AVG(amount) > 500;
```

| city | number_orders | average_amount |
|-------|---------------|----------------|
| Milan | 25 | 675.50 |
| Rome | 18 | 820.30 |

Multiple conditions in HAVING

The HAVING Clause

```
SELECT  
brand,  
COUNT(*) AS models,  
MIN(price) AS min_price,  
MAX(price) AS max_price  
FROM automobiles  
GROUP BY brand  
HAVING MAX(price) - MIN(price) > 20000;
```

| brand | models | min_price | max_price |
|----------|--------|-----------|-----------|
| BMW | 8 | 35000 | 85000 |
| Mercedes | 6 | 40000 | 95000 |

HAVING with **expressions** between aggregations

Limiting Result Tuples

LIMIT Clause

Allows you to limit the number of rows returned by the query.

Syntax:

```
SELECT columns  
FROM table  
WHERE condition  
ORDER BY column  
LIMIT number;
```

Common uses:

- Display only the first N results
- Implement pagination
- Optimize query performance

Limiting Result Tuples

```
SELECT name, surname, salary
FROM employees
ORDER BY salary DESC
LIMIT 5;
```

Top 5 highest salaries:

| name | surname | salary |
|----------|---------|--------|
| Anna | Neri | 45000 |
| Mario | Rossi | 42000 |
| Laura | Bianchi | 38000 |
| Giuseppe | Verdi | 35000 |
| Carlo | Gialli | 33000 |

Returns only the first 5 ordered rows

Limiting Result Tuples

Pagination with OFFSET

OFFSET allows you to skip a specified number of rows.

```
SELECT name, price
FROM products
ORDER BY price DESC
LIMIT 10 OFFSET 20;
```



Skip 20



Get 10



Others...

Useful for pagination: page 3 = LIMIT 10 OFFSET 20

Limiting Result Tuples

```
SELECT  
product ,  
sales ,  
revenue  
FROM sales_statistics  
ORDER BY revenue DESC  
LIMIT 3;
```

Top 3 products by revenue:

| product | sales | revenue |
|--------------|-------|----------|
| Laptop Pro | 234 | 234000 € |
| Smartphone X | 456 | 182400 € |
| Tablet Plus | 198 | 99000 € |

Best Practice

Always use ORDER BY with LIMIT for consistent results.

Limiting Result Tuples

Advantages of LIMIT

- **Performance:** reduces database load
- **Memory:** transfers less data
- **Usability:** improves user experience
- **Testing:** useful for testing queries on large datasets

```
-- Test query on large table  
SELECT * FROM event_logs  
WHERE date >= '2024-01-01'  
LIMIT 100;
```

